Neomundi. The AI thermometer.

A real-time stability signal your systems use to decide.

Your agents no longer wait for the response. They react to the signal.
During generation. Not after.

THE PROBLEM

AI responds. You assume the risk.

An incorrect legal response. An undetected drift. A hallucination nobody saw coming.

Today, you only find out after.

Every unmeasured response is an uncontrolled risk.

You no longer correct after. You decide during.

WHAT NEOMUNDI DOES

Neomundi injects a stability signal into your generation pipeline.

It sits between your application and the AI model. Each output receives a score. Each score is returned to your systems, which decide to stop, continue, or correct.

No code change. One single API call.

Input → score → decision. In real time.

WHAT YOU RECEIVE ON EACH GENERATION

- **Stability score** (0 → 1) — can it continue?
- **Decision** (ALLOW / FLAG / BLOCK) — actionable by your systems
- **Drift signal** (timestamped) — error anticipation

Measured signal. Immediate decision. Risk under control.

TWO OPERATIONAL MODES

OBS Mode : Observation & Score	GOV Mode : Active Governance
Measures and logs every response. No intervention. No infrastructure change. Available now.	Intercepts before emission. ALLOW, FLAG or BLOCK at the speed of generation. Available as pilot.

SIGNAL → ACTION → IMPACT

	Signal	Action	Impact
Agents	Live stability	Stop / Reroute	Reliable agents
Compliance	Drift detected	Proof generated	Auditable output
Fine-tuning	Gap measured	Clean dataset	Improved model
SLA / Infra	Quality measured	SLA documented	Continuity proved

VALIDATED RESULTS

79.5%	91%	26	<2s
Detection (TruthfulQA)	Precision <9% false alarms	Risk categories	Added latency

EU AI ACT

You are not preparing for compliance. You are already compliant, by architecture.

Every response scored, traced, auditable. By design.

POSITIONING

Portkey routes. LangSmith analyses. Helicone logs. Before or after. Never during.

Neomundi does not judge truth. It detects when your AI drifts at the speed of generation.

Not a firewall. Not a guardrail.

A new category: decision infrastructure based on runtime signal.

USE CASE: LEGAL PLATFORM

Before: the lawyer reviews and assumes the risk.

After: hallucination detected before the lawyer sees the response. Traceable score. Admissible proof. Liability is not eliminated it is measured and assigned.

You are no longer simply deploying AI. You are deploying AI with proof that you governed it.

Plug in NeoMundi. First signal in minutes.

neomundi.io **contact@neomundi.io**

Reliability is no longer assumed. It is continuously measured.

© 2026 NeoMundi Lab