NeoMundi

Executive Brief 2026

Every uncontrolled AI response is a risk you own alone.

The problem

A flawed legal response. A risky client recommendation. A response that drifts without anyone noticing.

The AI responds. You are liable.

Today, you only find out after.

What NeoMundi does

The NeoMundi API sits between your application and the AI model.

NeoMundi intervenes at the only moment that matters : **while the AI is generating.**

Every output receives a stability score. Every score drives a decision, in milliseconds, before the response reaches anyone.

No code changes. Deployment in minutes.

Stability score : 0 = unstable 1 = stable. Real time. Automatic. Provable.

Without NeoMundi	With NeoMundi
x Audited after the incident	v Flagged before output
x Drift goes undetected	v Drift detected in real time
x Reliability assumed	v Reliability measured and proven
x No proof in case of dispute	v Every response traceable, auditable, defensible

EU AI Act : auditability by design

You do not prepare for compliance.

You are already compliant, by architecture.

Every response scored, traced, auditable. By design.

Not a firewall. Not a guardrail.

The entire market observes or constrains. NeoMundi decides.

Portkey transports. LangSmith analyzes. Helicone logs. They act before or after. Never during.

NeoMundi is a runtime decision layer : **a new category.**

Validated results

79.5% detection of unstable responses before delivery (TruthfulQA, 790 questions)

91% precision : less than 9% false alarms

26 risk categories covered

< 2s average latency per response

Full results in the Technical White Paper.

Two operational modes

OBS Mode : Observation and Score

Measures and logs every response. No intervention. No infrastructure change. Available now.

GOV Mode : Active Governance Available in pilot.

Fewer tokens for a better response : the system knows when to stop, as soon as the useful information is delivered.

Use case : AI-assisted legal platform

A law firm deploys an LLM to pre-draft client responses.

Before :
the lawyer reviews and takes the risk.

After with NeoMundi :
the risk is detected before the lawyer even sees the response.

>> Hallucination above threshold, flagged automatically
>> Every response carries a score, traceable, auditable, defensible
>> Liability is not eliminated, it is measured and assigned

You are no longer just deploying AI.
You are deploying AI with proof that you governed it.

Connect NeoMundi. Get your first score in minutes.

neomundi.io contact@neomundi.io

Reliability is no longer assumed. It is continuously measured. 2026 NeoMundi Lab